

```
In [1]: import pandas as pd
import numpy as np
```

```
In [2]: df= pd.read_csv("C:\\Users\\ASHISH TIWARI\\Downloads\\Churn_Modelling.csv")
df
```

```
Out[2]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProduct
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	
3	4	15701354	Boni	699	France	Female	39	1	0.00	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	
...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	

10000 rows × 14 columns

```
In [3]: df.head()
```

```
Out[3]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	H
0	1	15634602	Hargrave	619	France	Female	42	2	0.00		1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86		1
2	3	15619304	Onio	502	France	Female	42	8	159660.80		3
3	4	15701354	Boni	699	France	Female	39	1	0.00		2
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82		1

```
In [4]: df.tail()
```

```
Out[4]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProduct
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	

```
In [5]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber             10000 non-null  int64
1   CustomerId            10000 non-null  int64
2   Surname               10000 non-null  object
3   CreditScore           10000 non-null  int64
4   Geography             10000 non-null  object
5   Gender                10000 non-null  object
6   Age                   10000 non-null  int64
7   Tenure                10000 non-null  int64
8   Balance               10000 non-null  float64
9   NumOfProducts        10000 non-null  int64
10  HasCrCard             10000 non-null  int64
11  IsActiveMember        10000 non-null  int64
12  EstimatedSalary       10000 non-null  float64
13  Exited                10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

```

In [6]: `df.dtypes`

```

Out[6]: RowNumber      int64
CustomerId    int64
Surname       object
CreditScore   int64
Geography     object
Gender        object
Age           int64
Tenure        int64
Balance       float64
NumOfProducts int64
HasCrCard     int64
IsActiveMember int64
EstimatedSalary float64
Exited        int64
dtype: object

```

In [7]: `df.describe(include='all')`

```

Out[7]:

```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	I
<b>count</b>	10000.00000	1.000000e+04	10000	10000.000000	10000	10000	10000.000000	10000.000000	10000
<b>unique</b>	NaN	NaN	2932	NaN	3	2	NaN	NaN	
<b>top</b>	NaN	NaN	Smith	NaN	France	Male	NaN	NaN	
<b>freq</b>	NaN	NaN	32	NaN	5014	5457	NaN	NaN	
<b>mean</b>	5000.50000	1.569094e+07	NaN	650.528800	NaN	NaN	38.921800	5.012800	76485
<b>std</b>	2886.89568	7.193619e+04	NaN	96.653299	NaN	NaN	10.487806	2.892174	62397
<b>min</b>	1.00000	1.556570e+07	NaN	350.000000	NaN	NaN	18.000000	0.000000	0
<b>25%</b>	2500.75000	1.562853e+07	NaN	584.000000	NaN	NaN	32.000000	3.000000	0
<b>50%</b>	5000.50000	1.569074e+07	NaN	652.000000	NaN	NaN	37.000000	5.000000	97198
<b>75%</b>	7500.25000	1.575323e+07	NaN	718.000000	NaN	NaN	44.000000	7.000000	127644
<b>max</b>	10000.00000	1.581569e+07	NaN	850.000000	NaN	NaN	92.000000	10.000000	250898

In [8]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber             10000 non-null  int64
1   CustomerId            10000 non-null  int64
2   Surname               10000 non-null  object
3   CreditScore           10000 non-null  int64
4   Geography             10000 non-null  object
5   Gender                10000 non-null  object
6   Age                   10000 non-null  int64
7   Tenure                10000 non-null  int64
8   Balance               10000 non-null  float64
9   NumOfProducts         10000 non-null  int64
10  HasCrCard             10000 non-null  int64
11  IsActiveMember        10000 non-null  int64
12  EstimatedSalary        10000 non-null  float64
13  Exited                10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

```

In [9]: `df.isnull().sum()`

```

Out[9]: RowNumber      0
        CustomerId    0
        Surname       0
        CreditScore   0
        Geography     0
        Gender        0
        Age           0
        Tenure        0
        Balance       0
        NumOfProducts 0
        HasCrCard     0
        IsActiveMember 0
        EstimatedSalary 0
        Exited        0
dtype: int64

```

In [10]: `df.columns`

```

Out[10]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
                'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
                'IsActiveMember', 'EstimatedSalary', 'Exited'],
               dtype='object')

```

In [11]: `df = df.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1)`

In [12]: `df.head()`

```

Out[12]:
   CreditScore  Geography  Gender  Age  Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember  EstimatedSalary
0          619    France  Female   42     2     0.00             1           1             1             101
1          608    Spain  Female   41     1  83807.86             1           0             1             112
2          502    France  Female   42     8  159660.80             3           1             0             113
3          699    France  Female   39     1     0.00             2           0             0              93
4          850    Spain  Female   43     2  125510.82             1           1             1              79

```

In [13]: `df.columns`

```

Out[13]: Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
                'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary',
                'Exited'],
               dtype='object')

```

In [14]: `df['Geography'].unique()`

```

Out[14]: array(['France', 'Spain', 'Germany'], dtype=object)

```

```
In [15]: df = pd.get_dummies(df, drop_first=True)
df.head()
```

```
Out[15]:
```


	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited	Geogr
0	619	42	2	0.00	1	1	1	101348.88	1	
1	608	41	1	83807.86	1	0	1	112542.58	0	
2	502	42	8	159660.80	3	1	0	113931.57	1	
3	699	39	1	0.00	2	0	0	93826.63	0	
4	850	43	2	125510.82	1	1	1	79084.10	0	

```
In [16]: df['Exited'].value_counts()
```

```
Out[16]: Exited
0      7963
1      2037
Name: count, dtype: int64
```

```
In [24]: import seaborn as sns
sns.countplot(x='Exited', hue='Exited', data=df)
```

```
Out[24]: <Axes: xlabel='Exited', ylabel='count'>
```



```
In [29]: X = df.drop('Exited', axis = 1)
Y = df['Exited']
```

```
In [30]: from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.20, random_state= 51, stratify=Y)
```

```
In [31]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
X_train
```

```
Out[31]: array([[ -0.78860057, -1.03742686,  0.6965047 , ..., -0.57619557,
          1.73726261, -1.09582175],
          [ -0.68522026, -0.56111692, -0.68966051, ...,  1.73552185,
          -0.57561821,  0.91255717],
          [ -1.03671334, -0.08480698,  1.3895873 , ..., -0.57619557,
          -0.57561821,  0.91255717],
          ...,
          [  1.34103398, -0.46585493,  0.0034221 , ...,  1.73552185,
          -0.57561821,  0.91255717],
          [ -0.20967079,  0.96307488, -1.38274311, ..., -0.57619557,
          1.73726261, -1.09582175],
          [  1.12393531,  0.010455 ,  0.3499634 , ..., -0.57619557,
          -0.57561821, -1.09582175]])
```

```
In [32]: from sklearn.linear_model import LogisticRegression
log = LogisticRegression()
log.fit(X_train,Y_train)
```

```
Out[32]: ▾ LogisticRegression
LogisticRegression()
```

```
In [33]: Y_pred1 = log.predict(X_test)
```

```
In [34]: from sklearn.metrics import accuracy_score
accuracy_score(Y_test,Y_pred1)
```

```
Out[34]: 0.81
```

```
In [35]: from sklearn.neighbors import KNeighborsClassifier
knn= KNeighborsClassifier()
knn.fit(X_train, Y_train)
KNeighborsClassifier()
```

```
Out[35]: ▾ KNeighborsClassifier
KNeighborsClassifier()
```

```
In [36]: Y_pred2 = knn.predict(X_test)
```

```
In [37]: accuracy_score(Y_test, Y_pred2)
```

```
Out[37]: 0.827
```

```
In [38]: from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier()
rf.fit(X_train, Y_train)
```

```
Out[38]: ▾ RandomForestClassifier
RandomForestClassifier()
```

```
In [39]: Y_pred3 = rf.predict(X_test)
```

```
In [40]: accuracy_score(Y_test, Y_pred3)
```

```
Out[40]: 0.8635
```

```
In [41]: from sklearn.ensemble import GradientBoostingClassifier
gradient_booster = GradientBoostingClassifier()
gradient_booster.fit(X_train,Y_train)
```

```
Out[41]: ▾ GradientBoostingClassifier
GradientBoostingClassifier()
```

```
In [42]: Y_pred4 = gradient_booster.predict(X_test)
```

In [43]: `accuracy_score(Y_test, Y_pred4)`

Out[43]: **0.8615**

In [ ]: